

Object Classification with SOTA Models


MSc. Nguyen Thanh An
VietAI Teaching Team

Content

- Object classification problem
- SOTA CNN models
 - VGG
 - Res-Net
 - Mobile-Net
 - Efficient-Net
- Attention Mechanism
- Vision Transformer
- Swin Transformer

Object Classification Problem

- Classification
 - Input: an image
 - Output: category/ID of the image
- Examples
 - Face recognition
 - Skin lesion classification
 - Fish classification
 - etc.



<https://isysrg.com/2020/06/17/deep-ensemble-learning-for-skin-lesions-classification-with-convolutional-neural-network/>

Image Representation

- Images are represented as matrices of pixels in computer.
 - Grayscale image → pixel is a single integer
 - Color image → pixel is a tuple of integers (Red, Green, Blue)
 - Other color spaces: HSV, CMYK, etc.
- Images are usually transformed to meet the structure of Machine Learning model input.
 - Flatten to a vector
 - Resize
 - Color space transformation



<https://en.wikipedia.org/wiki/Lenna>




https://www.researchgate.net/figure/8-bit-256-x-256-Grayscale-Lena-Image_fig1_3935609

Classification Approaches

- Handcrafted features
 - Extract discriminative visual pattern in the image → SIFT, SURF, BLOB, HOG, etc.
 - Represent the image as a set of extracted features
 - Apply Machine Learning models to classify these features → SVM, Random Forest, etc.
- Feature-learning
 - Utilize raw images → matrices of pixels
 - Feed the image to a neural network → convolutional neural network
 - Obtain the prediction


Feature-Learning Approach

- Deep Learning based models, especially convolutional neural network, are common applied.
- Deep Learning pipeline for object classification



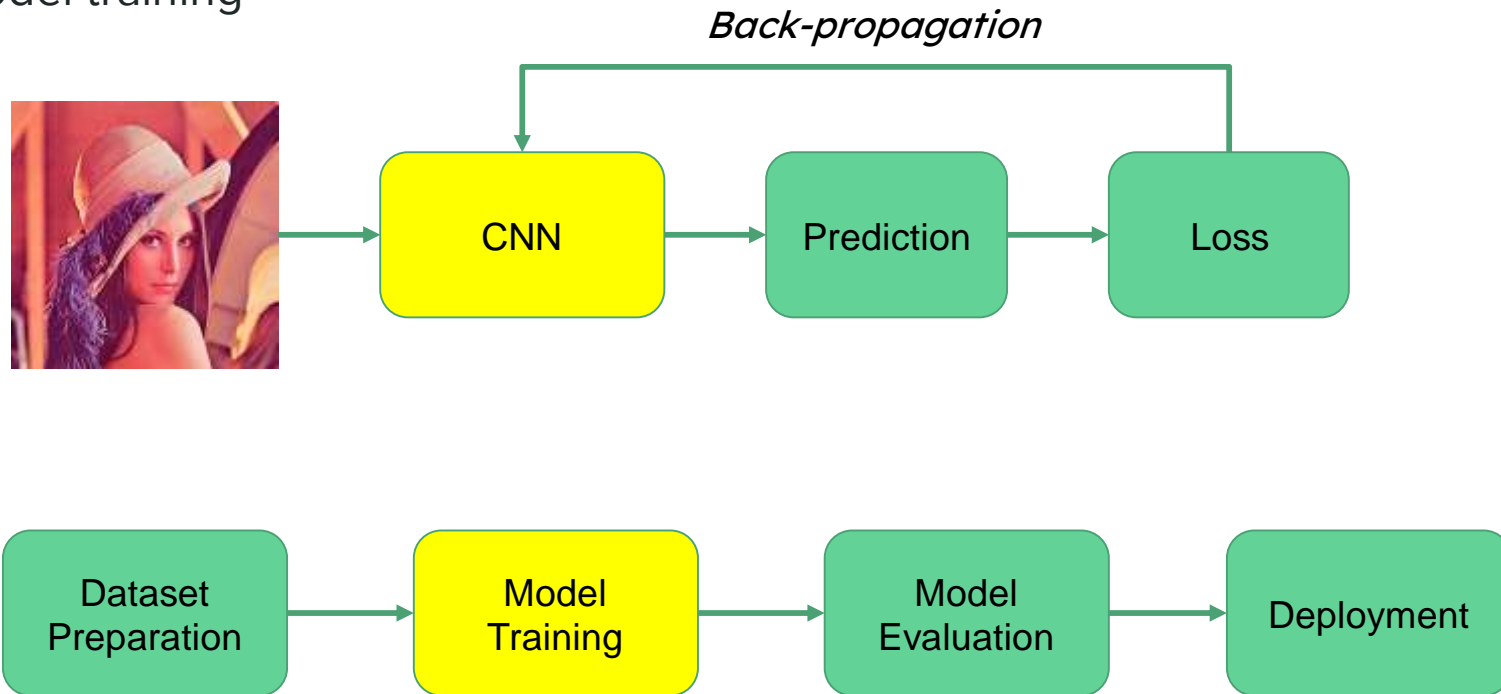
Feature-Learning Approach

- Dataset Preparation:
 - Samples = {image, label}
 - Preprocessing: cropping, normalization, resizing, etc.



Feature-Learning Approach


- Model training



Feature-Learning Approach

- Model Evaluation: common metrics for classification

- Precision
- Recall
- F1-score




How many retrieved items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are retrieved?


$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

<https://en.wikipedia.org/wiki/F-score>



Feature-Learning Approach

- Deployment:
 - Implement the pretrained model (with pretrained weights)
 - Integrate into a software.
 - Programming languages and frameworks are carefully considered to optimize performance.



Convolutional Neural Network

- CNN → end-to-end model for image classification
- General architecture
 - Input layer
 - Convolutional layer + ReLU → Max-Pooling layer
 - ...
 - Convolutional layer + ReLU → Max-Pooling layer
 - Flattenning
 - Dense layer + Sigmoid + Dropout
 - ...
 - Dense layer + Sigmoid + Dropout
 - Softmax layer
 - Output

Feature extraction phase

Classification phase


Output $[H'', W'', N]$



History of CNNs

- History of CNNs with typical SOTA models
 - 2014 VGG → <https://arxiv.org/abs/1409.1556>
 - 2015 Res-Net → <https://arxiv.org/abs/1512.03385>
 - 2017 Mobile-Net → <https://arxiv.org/abs/1704.04861>
 - 2019 Efficient-Net → <https://arxiv.org/abs/1905.11946>


VGG-16



- Depth of a convolutional neural network takes effect of its classification accuracy.
- Shadow networks (with fewer layers) have a smaller number of parameters → underfitting
- Deep networks (with more layers) have a larger number of parameters → overfitting
- VGG's authors: replace a conv. layer with a large kernel size, i.e. 5×5 , by a stack of ones with smaller kernel size, i.e. 3×3 → maintain the same effect, but smaller number of parameters.



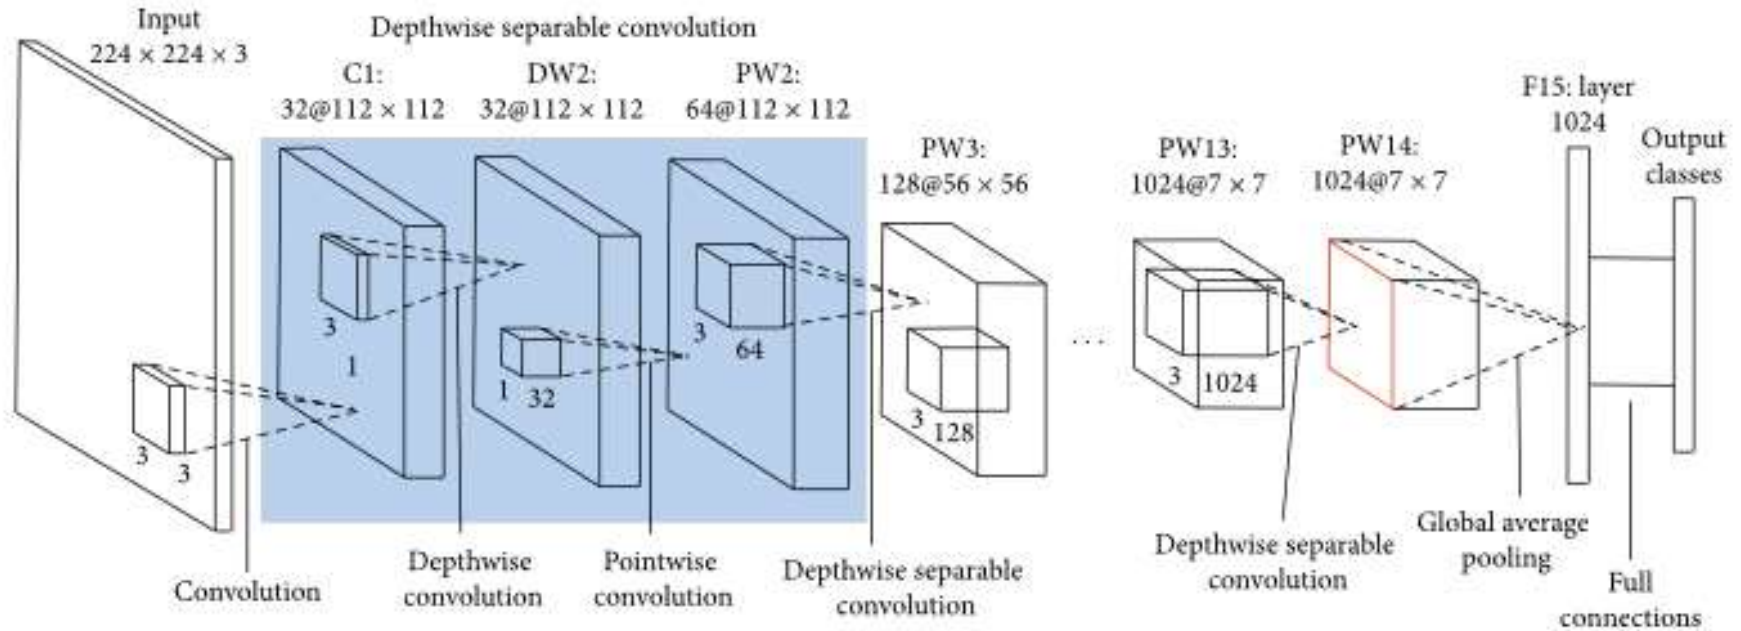
Residual Block




<https://www.geeksforgeeks.org/residual-networks-resnet-deep-learning/?ref=lbp>

- If any layer hurts the performance of architecture, it will be skipped by regularization.
- Train a very deep neural network without vanishing/exploding gradient.
- The idea of residual blocks can be inherited and customized in later models.


Mobile-Net




<https://medium.com/analytics-vidhya/image-classification-with-mobilenet-cc6fbb2cd470>



(a) Standard Convolution Filters



(b) Depthwise Convolutional Filters

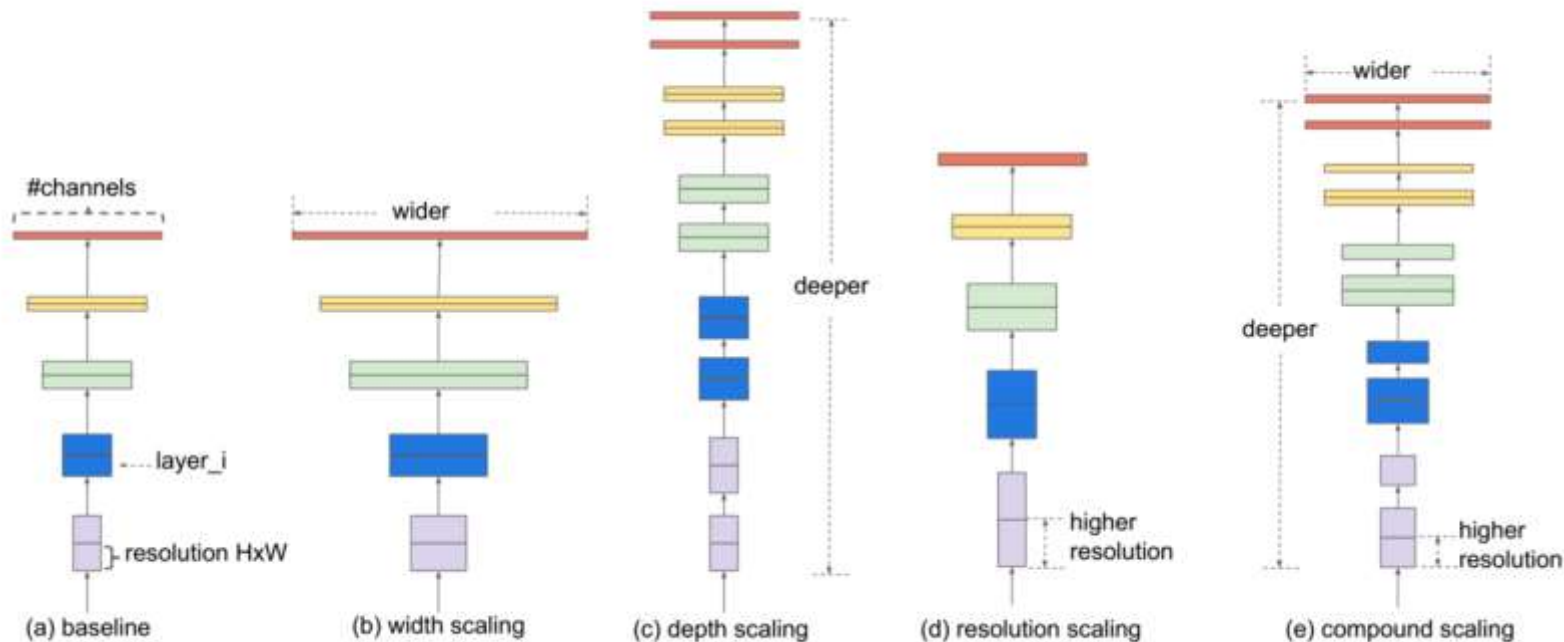


(c) 1×1 Convolutional Filters called Pointwise Convolution in the context of Depthwise Separable Convolution

<https://arxiv.org/pdf/1704.04861.pdf>

- Streamlined architecture that uses depth-wise separable convolutions to build light weight deep neural networks
- Target to mobile and embed devices.
- Depth-wise convolutional filters reduce computation → lightweight model.

Efficient-Net



<https://arxiv.org/pdf/1905.11946.pdf>

- EfficientNet's authors propose a systematic method for model scaling that helps improve accuracy and resource efficiency.
- Previous models are scaled randomly based on three major concepts: depth, width, and resolution → manually tuning
- Scales each dimension with a certain fixed set of scaling coefficients

$$\text{depth: } d = \alpha^\phi$$

$$\text{width: } w = \beta^\phi$$

$$\text{resolution: } r = \gamma^\phi$$

$$\text{s.t. } \alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$$

$$\alpha \geq 1, \beta \geq 1, \gamma \geq 1$$

Table 1. EfficientNet-B0 baseline network – Each row describes a stage i with \hat{L}_i layers, with input resolution $\langle \hat{H}_i, \hat{W}_i \rangle$ and output channels \hat{C}_i . Notations are adopted from equation 2.

Stage i	Operator $\hat{\mathcal{F}}_i$	Resolution $\hat{H}_i \times \hat{W}_i$	#Channels \hat{C}_i	#Layers \hat{L}_i
1	Conv3x3	224×224	32	1
2	MBConv1, k3x3	112×112	16	1
3	MBConv6, k3x3	112×112	24	2
4	MBConv6, k5x5	56×56	40	2
5	MBConv6, k3x3	28×28	80	3
6	MBConv6, k5x5	14×14	112	3
7	MBConv6, k5x5	14×14	192	4
8	MBConv6, k3x3	7×7	320	1
9	Conv1x1 & Pooling & FC	7×7	1280	1

Efficient-Net

Model	Top-1 Acc.	Top-5 Acc.	#Params	Ratio-to-EfficientNet	#FLOPs	Ratio-to-EfficientNet
EfficientNet-B0	77.1%	93.3%	5.3M	1x	0.39B	1x
ResNet-50 (He et al., 2016)	76.0%	93.0%	26M	4.9x	4.1B	11x
DenseNet-169 (Huang et al., 2017)	76.2%	93.2%	14M	2.6x	3.5B	8.9x
EfficientNet-B1	79.1%	94.4%	7.8M	1x	0.70B	1x
ResNet-152 (He et al., 2016)	77.8%	93.8%	60M	7.6x	11B	16x
DenseNet-264 (Huang et al., 2017)	77.9%	93.9%	34M	4.3x	6.0B	8.6x
Inception-v3 (Szegedy et al., 2016)	78.8%	94.4%	24M	3.0x	5.7B	8.1x
Xception (Chollet, 2017)	79.0%	94.5%	23M	3.0x	8.4B	12x
EfficientNet-B2	80.1%	94.9%	9.2M	1x	1.0B	1x
Inception-v4 (Szegedy et al., 2017)	80.0%	95.0%	48M	5.2x	13B	13x
Inception-resnet-v2 (Szegedy et al., 2017)	80.1%	95.1%	56M	6.1x	13B	13x
EfficientNet-B3	81.6%	95.7%	12M	1x	1.8B	1x
ResNeXt-101 (Xie et al., 2017)	80.9%	95.6%	84M	7.0x	32B	18x
PolyNet (Zhang et al., 2017)	81.3%	95.8%	92M	7.7x	35B	19x
EfficientNet-B4	82.9%	96.4%	19M	1x	4.2B	1x
SENet (Hu et al., 2018)	82.7%	96.2%	146M	7.7x	42B	10x
NASNet-A (Zoph et al., 2018)	82.7%	96.2%	89M	4.7x	24B	5.7x
AmoebaNet-A (Real et al., 2019)	82.8%	96.1%	87M	4.6x	23B	5.5x
PNASNet (Liu et al., 2018)	82.9%	96.2%	86M	4.5x	23B	6.0x
EfficientNet-B5	83.6%	96.7%	30M	1x	9.9B	1x
AmoebaNet-C (Cubuk et al., 2019)	83.5%	96.5%	155M	5.2x	41B	4.1x
EfficientNet-B6	84.0%	96.8%	43M	1x	19B	1x
EfficientNet-B7	84.3%	97.0%	66M	1x	37B	1x
GPipe (Huang et al., 2018)	84.3%	97.0%	557M	8.4x	-	-

<https://arxiv.org/pdf/1905.11946.pdf>

Attention Mechanism

- Attention mechanism
 - Enhance important parts and
 - Fade out non-relevant information
- Simple example in NLP:
 - Self-attention compares every word in a sentence to each other
 - Reweighting the embeddings of each word to include contextual relevance
- E.g. : bank of the river
 - “River” changes the contextual meaning of “bank”
 - If we do not see “river”, then misunderstanding

Basic Attention Module


- Basic structure of an attention module
 - Tensor x_1 \rightarrow generates a “key” and a “value”
 - Tensor x_2 \rightarrow generates a “query”
 - Output is the weighted sum of “value”
 - Weights are computed by compatibility function of “query” and the corresponding “key”

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Multi-head Attention

- Multi-head attention:
 - runs an attention module several times and concatenate outputs.
 - jointly attend to information from different representation subspaces at different positions (not able for single-head attention)

<https://paperswithcode.com/method/multi-head-attention>



Convolutional Block Attention

- Convolutional Block Attention: emphasizes meaningful features along
 - the channel and
 - spatial axes
- Applicable at every convolutional block

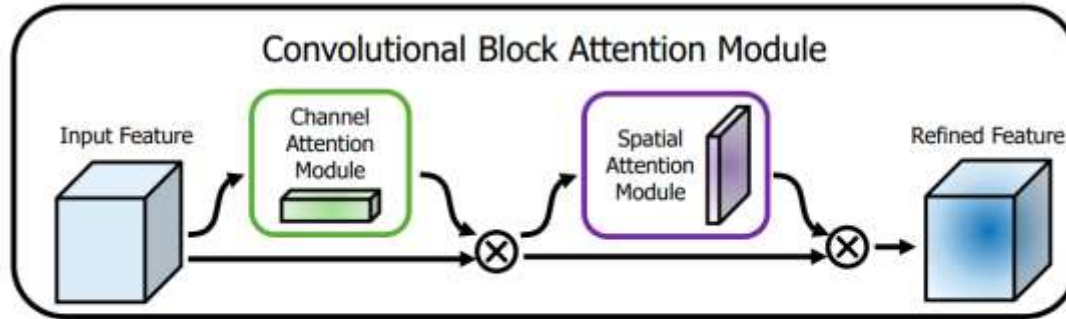


Fig. 1: **The overview of CBAM.** The module has two sequential sub-modules: *channel* and *spatial*. The intermediate feature map is adaptively refined through our module (CBAM) at every convolutional block of deep networks.

Convolutional Block Attention

- Channel Attention Module (CAM):
 - input tensor \rightarrow 2 vectors ($c \times 1 \times 1$) (Global Average Pooling and Global Max Pooling)
 - output \rightarrow a fully connected layer + ReLU

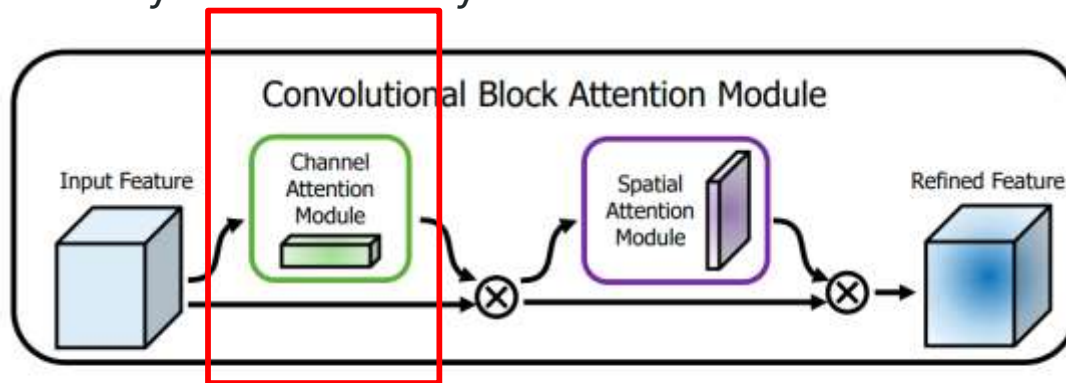
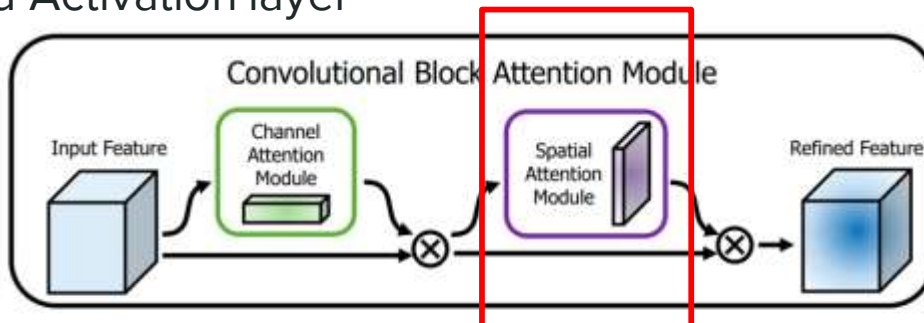


Fig. 1: **The overview of CBAM.** The module has two sequential sub-modules: *channel* and *spatial*. The intermediate feature map is adaptively refined through our module (CBAM) at every convolutional block of deep networks.

Convolutional Block Attention

- Spatial Attention Module (SAM):
 - Channel Pool: apply Max Pooling and Average Pooling across the channels
 - $(c \times h \times w) \rightarrow (2 \times h \times w)$.
 - Conv. layer + BatchNorm + ReLU: $(1 \times h \times w)$
 - Sigmoid Activation layer



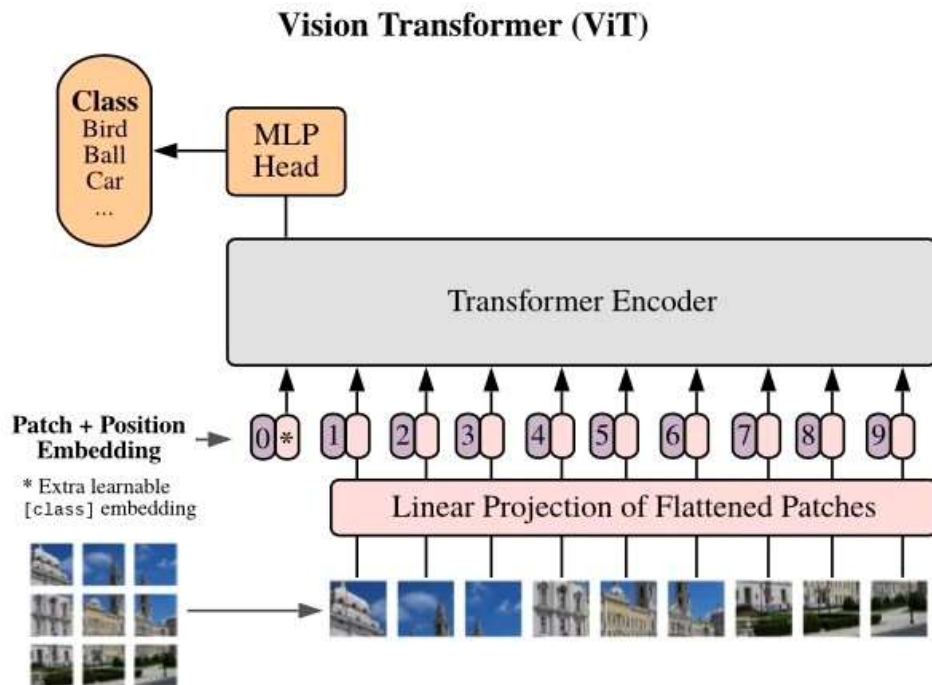
<https://arxiv.org/pdf/1807.06521.pdf>

Fig. 1: **The overview of CBAM.** The module has two sequential sub-modules: *channel* and *spatial*. The intermediate feature map is adaptively refined through our module (CBAM) at every convolutional block of deep networks.


Vision Transformer

- Vision Transformer is a recent advanced technique without reliance on convolutional neural networks.
 - Empirical results surpass CNNs'
 - Lower computational cost
- Main idea:
 - Divide images into patches encoded with position information
 - Handle encodings using attention mechanism
 - Classify output tensors using MLP.

Vision Transformer



Transformer Encoder

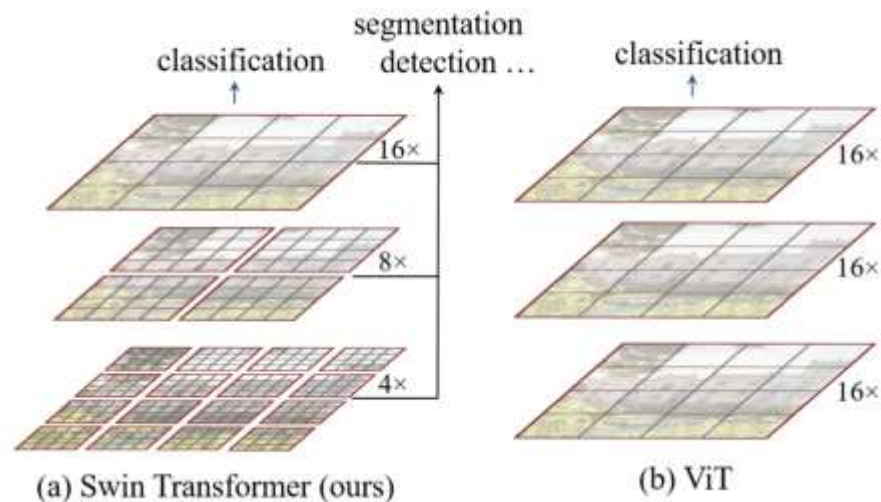


Procedure of Vision Transformer

- Divide the input image into patches of a fixed size
- Flatten patches
- Generate feature embeddings with lower dimensionality from patches
- Attach order (position) of patches in feature embeddings
- Feed embeddings to a transformer encoder
- Train and evaluate the model in the whole dataset
- Tune model hyperparameters regarding to individual problems/cases

Swin Transformer

- Image resolution is a challenge of ViT, which is different from texts.
- Swin transformer is a hierarchical transformer whose representation is computed with Shifted windows.
- Shifted window based self-attention → reduce computational cost.



<https://arxiv.org/pdf/2103.14030.pdf>

Swin Transformer

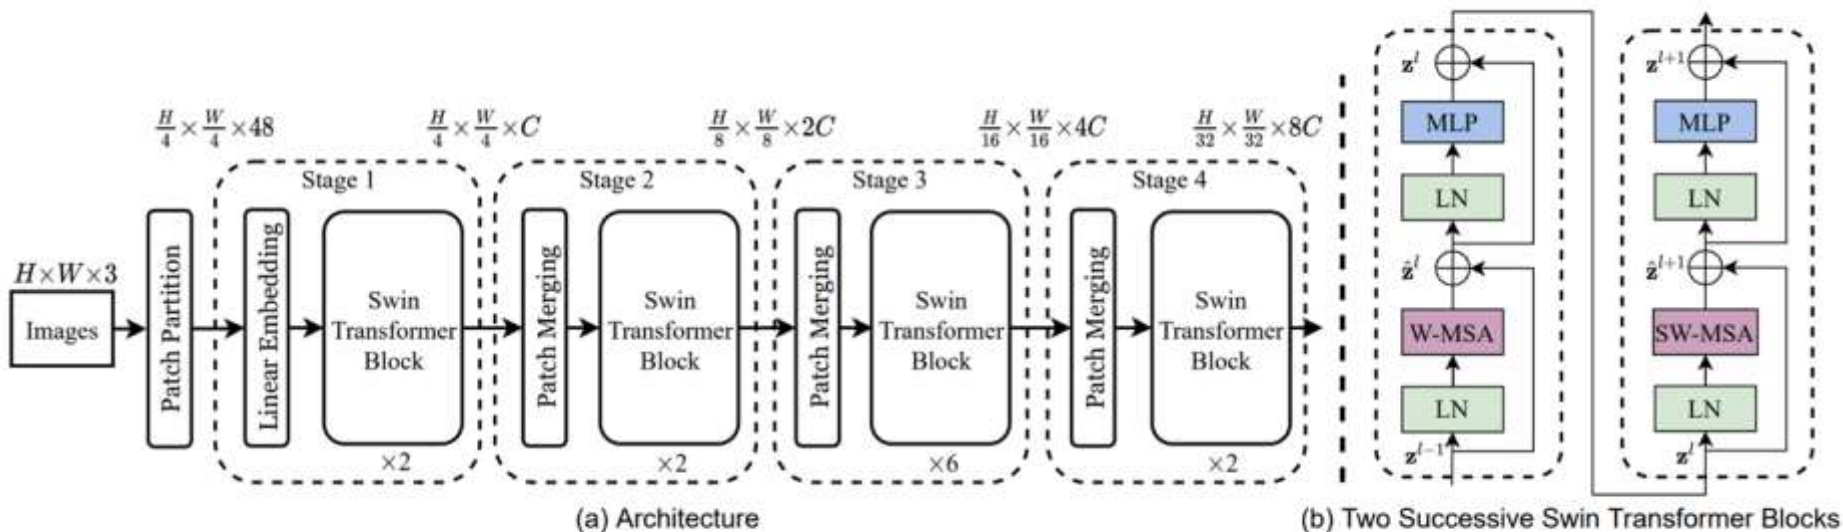


Figure 3. (a) The architecture of a Swin Transformer (Swin-T); (b) two successive Swin Transformer Blocks (notation presented with Eq. (3)). W-MSA and SW-MSA are multi-head self attention modules with regular and shifted windowing configurations, respectively.

Swin Transformer

- Images \rightarrow non-overlapping patches
- Linear projection(patch) \rightarrow features
- Swin transformer blocks
- Patch merging (neighbor patches of each group 2×2)
 \rightarrow patch size increases after each “stage”

Swin Transformer

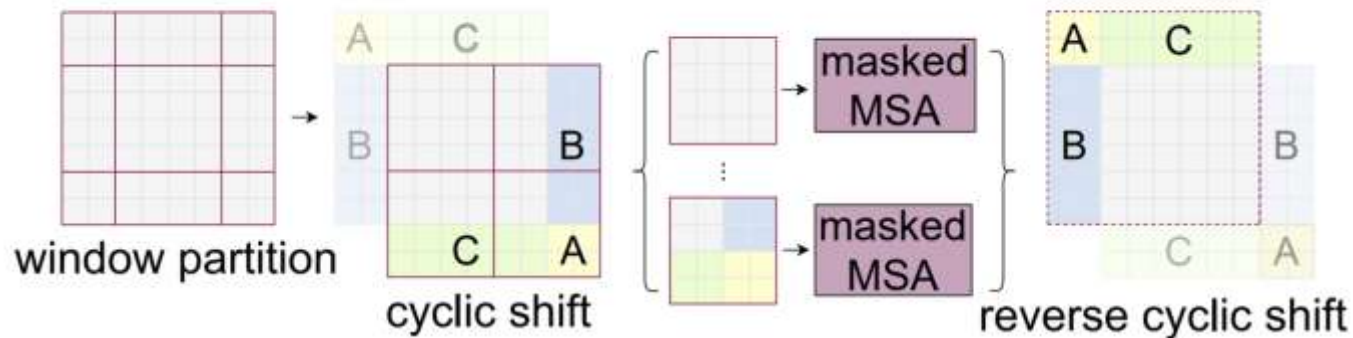


Figure 4. Illustration of an efficient batch computation approach for self-attention in shifted window partitioning.

Swin Transformer

- Shift the windows cyclically → windows with small size at borders and corners.
- Repeat the image periodically instead of zero-padding.
- “Padded regions” are handled using masks → limit self-attention computation to within non-adjacent sub-windows.

